



WikiText-103, 192d / 12 layers



Parameter budget



matched: context 8196, 32,784 tokens/optimizer step, 781 steps/epoch, 100 epochs, dropout 0.1, muon_hybrid
 parameters are the trained model; aux is training-only and discarded at inference